

Harsha Vardhan Vinjamuri

SUMMARY

Results-driven **Data Engineer with 5+ years of experience** designing enterprise-scale data pipelines, real-time streaming architectures, and cloud integration solutions across **financial services, healthcare, aerospace, and manufacturing** domains. Proficient in **Palantir Foundry, Apache Spark, PySpark, Kafka, and MuleSoft**, delivering high-performance ETL/ELT workflows and ML feature engineering pipelines processing **billions of records daily** across **AWS and Azure** environments. Experienced integrating complex enterprise ecosystems including **SAP S/4HANA, MES, PLM, and IoT systems**, with strong knowledge of data governance and security compliance (**HIPAA, SOC 2, FedRAMP, ITAR**). Skilled in DevOps automation using **Docker, Kubernetes, and CI/CD pipelines**, consistently delivering scalable, reliable data solutions in **Agile** environments.

SKILLS

- Programming & Scripting:** Python, SQL, PySpark, Scala, Shell Scripting (Unix)
- Big Data & Distributed Computing:** Apache Spark, PySpark, Spark SQL, Apache Hadoop, Apache Hive, Apache Kafka, Amazon EMR
- Data Engineering & ETL Tools:** dbt (Data Build Tool), Apache Airflow, MuleSoft Anypoint, Foundry Pipeline Builder, Azure Data Factory, AWS Glue, SQL Server Integration Services (SSIS)
- Cloud Platforms & Services:** AWS (EC2, S3, Glue, Redshift, RDS, DynamoDB), Azure (Databricks, Synapse Analytics, Data Lake, Service Bus, API Management, Azure DevOps), Palantir Foundry, Palantir AIP
- Databases & Data Warehousing:** Snowflake, SQL Server, PostgreSQL, MySQL, Oracle, MongoDB, Teradata
- Streaming & Real-Time Data:** Apache Kafka, RabbitMQ, Amazon Kinesis, Azure Event Hubs, Spark Streaming, MQTT, OPC UA, IoT Data Streams
- Machine Learning & AI:** Palantir AIP, PySpark ML, Scikit-learn, Azure ML, TensorFlow, PyTorch, LangChain, RAG, Pandas, NumPy, SciPy, Feature Engineering, Model Deployment, Hyperparameter Tuning
- DevOps & Infrastructure as Code:** Git, GitLab CI/CD, Jenkins, Docker, Kubernetes (AKS), Terraform, Azure DevOps
- Data Visualization & BI:** Power BI, Tableau, Microsoft Excel, QlikView, Seaborn, Matplotlib, SQL Server Reporting Services (SSRS)
- Methodologies:** Agile, Scrum, Waterfall, Software Development Life Cycle (SDLC)

EXPERIENCE

- Palantir | NY | Data Engineer** May 2024 – Present
- Designed and implemented **enterprise data integration pipelines on Palantir Foundry** using **Apache Spark, PySpark, SQL and Python** ingesting **real-time trading feeds, transactional records, and market data streams** to support **fraud detection, risk scoring and compliance analytics** processing **\$2B+ in daily transaction volume**.
 - Built **ML feature engineering and feature store pipelines** using **Palantir AIP, PySpark and Spark SQL** transforming raw transactional and behavioral datasets into **model-ready feature sets**, improving **fraud model precision by 23%** and reducing **false positives by 31%**.
 - Optimized **Spark-based ETL/ELT transformation jobs and Foundry Pipeline Builder workflows** through **incremental processing, partition pruning, caching and cluster tuning**, reducing data latency from **daily to 15-minute refresh cycles** and lowering **compute costs by 34%**.
 - Supported **healthcare data integration initiatives**, ingesting **40M+ de-identified EHR and clinical datasets** from **HL7/FHIR sources, REST APIs, and cloud storage** into **Foundry ontology models** for patient outcome analytics while ensuring full **HIPAA compliance**.
 - Automated deployment of Foundry transformation pipelines using **Apache Airflow, GitLab CI/CD, Docker and infrastructure-as-code practices**, standardizing build, test, and release workflows across client workstreams and reducing **deployment errors by 25%**.
 - Implemented **enterprise data governance, lineage, metadata management, and role-based access control (RBAC)** within Palantir Foundry, ensuring compliance with **SOC 2 Type II, HIPAA, and FedRAMP** mentored junior engineers on **Foundry development best practices and PySpark performance optimization**.
- Cyient | India | Data Engineer** January 2020 – May 2023
- Designed and implemented **enterprise data integration pipelines** using **MuleSoft, Azure Service Bus**, enabling **real-time data exchange** between **SAP S/4HANA, Solumina MES, Teamcenter PLM**, and **2,500+ IoT devices** across 8 aerospace manufacturing sites.
 - Built large-scale **event-driven data pipelines** using **Apache Kafka and RabbitMQ**, processing **450M+ streaming telemetry events** daily with **99.96% reliability** and sub-3-second latency to support **real-time analytics** and **operational data synchronization**.
 - Engineered **data transformation pipelines** and **protocol integration layers** (OPC UA, MQTT, Modbus, EDI X12, REST/SOAP APIs) to connect legacy industrial systems with modern **cloud data platforms**, ensuring **99%+ data accuracy and consistency**.
 - Implemented **CI/CD pipelines, Docker containers and Kubernetes (AKS)** deployments using **Azure DevOps** improving release efficiency and maintaining **99.9% uptime** for production data services.
 - Applied **data security, API security and governance controls** using **OAuth 2.0, mutual TLS (mTLS)** and **API Management**, ensuring compliance with **ITAR, AS9100, and SOC 2** standards in regulated environments.

EDUCATION

Master of Information Systems, Cleveland State University, Cleveland, Ohio